

White Paper  
Emerging Compute Models

# Understanding Alternative Compute Models

Driven by ever-increasing pressure on a multitude of issues – including cost control, manageability, security, regulatory compliance, and business continuity – some IT managers are considering centralized compute models as an alternative to traditional distributed software deployment.

Thanks to recent advances in technology, IT has a range of centrally-managed compute models from which to choose. From Terminal Services to VDI, from blade PCs through application and OS streaming, each model has its advantages and trade-offs.

Rarely is a single compute model adequate to meet the needs of all users and all applications. Instead, the needs of each group of users, as well as IT requirements and existing infrastructure, must be carefully considered. In most cases, an enterprise solution will likely involve a combination of models.

Although thin terminals may be suitable hardware for certain purely server-based models, rich PC platforms based on technology such as Intel® Core™2 processor with vPro™ technology for desktops and Intel® Centrino® 2 with vPro™ technology for notebooks offer highly manageable and secure platforms on which to deploy a wide range of solutions.

This white paper describes the variety of compute models available, how each works, and the advantages and limitations of each model for key vectors such as security, manageability, and power consumption. Also identified are the major providers of software solutions and the end-point devices to deploy each model.



## Table of Contents

Issues to Consider .....	3
Table 1. Compute Model Comparison.....	3
Terminal Services .....	4
Summary .....	4
How it Works.....	4
Figure 1. Terminal Services Architecture .....	4
Table 2. Advantages of Terminal Services .....	5
Table 3. Limitations of Terminal Services.....	5
Virtual Desktop Infrastructure (VDI) .....	6
Summary .....	6
How it Works.....	6
Figure 2. VDI Architecture .....	6
Table 4. Advantages of VDI.....	7
Table 5. Limitations of VDI.....	7
Blade PCs .....	8
Summary .....	8
How it Works.....	8
Figure 3. One-to-One Blade Architecture .....	8
Figure 4. One-to-Many Blade Architecture.....	8
Table 6. Advantages of Blade PCs.....	9
Table 7. Limitations of Blade PCs .....	9
OS Image Streaming or Remote OS Boot .....	10
Summary .....	10
How it Works.....	10
Figure 5. OS Image Streaming Architecture.....	10
Figure 6. Remote OS Boot Architecture.....	10
Table 8. Advantages of OS Image Streaming and Remote OS Boot.....	11
Table 9. Limitations of OS Image Streaming and Remote OS Boot.....	11
Application Streaming or Application Virtualization.....	12
Summary .....	12
How it Works.....	12
Figure 7. Application Streaming or Application Virtualization Architecture .....	12
Table 10. Advantages of Application Streaming or Application Virtualization.....	13
Table 11. Limitations of Application Streaming or Application Virtualization .....	13
Virtual Containers .....	14
Summary .....	14
How it Works.....	14
Figure 8. Virtual Containers Architecture .....	14
Table 11. Advantages of Virtual Containers.....	15
Table 12. Limitations of Virtual Containers .....	15
Factors to Consider .....	16
Conclusion .....	16
Additional Resources.....	16

## Issues to Consider

Here are some considerations to keep in mind in evaluating centralized models for your organization.

Be sure to consider the application and business needs of the entire user base to understand which applications and data it makes sense to centralize versus installing locally.

It's helpful to segment users based on tasks performed and applications required. Individual users may need to access different applications using a mix of models. Rarely will a single model meet all the needs of a given group of users, let alone the needs of every user in the enterprise.

Decisions about which compute model to use often get intertwined with the client device on which it will be deployed. It's important to consider these topics separately.

For example, your business scenario may dictate server-based computing for a certain application. However, this server-based model, often referred to as "thin client," does not necessarily have to be deployed on a thin terminal.

A desktop or laptop PC may actually be a more appropriate device, depending on the user's total application and mobility needs. The user may also require locally executed applications. The client device should accommodate the most demanding requirements of target users.

For cost comparisons, it is suggested that the total cost of the complete solutions are compared under different compute models. Costs can just shift from the clients to the data center and vice versa, without actually reducing the total delivery cost. Also, infrastructure costs for equivalent processing and storage capacity can be significantly different depending on whether these are located in an endpoint or delivered from a data center.

**Table 1.** Compute Model Comparison

	<b>Terminal Services</b>	<b>Virtual Desktop Infrastructure (VDI)</b>	<b>Blade PCs</b>	<b>OS Image Streaming or Remote OS Boot</b>	<b>Application Streaming or Application Virtualization</b>	<b>Virtual Containers</b>
<b>Application Execution</b>	Server	Server	Server	Client	Client	Client
<b>Application Data Storage</b>	Server	Server	Server	Server	Client or Server	Client or Server
<b>Local Device Connect and Synch (Bar Code Reader, PDA, Phone, etc.)</b>	Limited	Limited	Yes	Yes	Yes	Limited
<b>Off-Network Mobile Option</b>	No	No	No	No	Yes	Yes
<b>Typical Clients</b>	Terminal, desktop PC, laptop PC	Terminal, desktop PC, laptop PC	Terminal, desktop PC	Desktop PC	Desktop PC, laptop PC	Desktop PC, laptop PC
<b>Largest Solution Providers</b>	Citrix,* Microsoft	Citrix, Microsoft, Red Hat, VMware	ClearCube, HP	Citrix, Dell, Lenovo	Citrix, Microsoft, Symantec, VMware	Citrix, <sup>a</sup> Microsoft, VMware

<sup>a</sup>Product announcement expected for 2009 release

# Server-Based Models

## Terminal Services

### Summary

Terminal Services is a time-tested and reliable server-based model dating back to mainframe computing. For supported software applications, this model offers strong security and manageability. However, users accustomed to the PC experience may find this model unsatisfying in terms of performance, customization, flexibility, and mobility.

Most large enterprises use Terminal Services for some applications, especially where security is essential and users are in a fixed location with constant network access. Bank tellers accessing the transaction system, call center workers entering orders, and healthcare professionals working with patient records are examples of where Terminal Services may be a good solution.

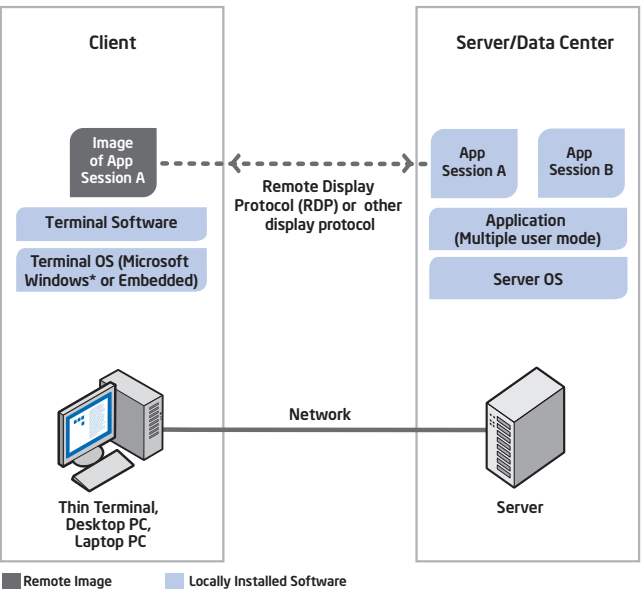
PCs are a good platform when Terminal Services is used only for a few key applications, while the rest are locally installed. Thin clients are appropriate only for an environment consisting of 100 percent

Terminal Services. And even then, replacing large numbers of existing PCs with new terminals can be costly. These costs should be closely evaluated to determine true return on investment (ROI). Many IT organizations choose to use existing or waterfalled PCs as Terminal Services' clients.

### How it Works

The client is merely a display and input device. All computation is done centrally on the server, and all data is stored in a data center. Nothing is executed or persistent on the client. Usually, Remote Display Protocol (RDP) or Independent Computing Architecture\* (ICA\*) is used to push an image of the server-based application to a terminal viewer on the client.

Figure 1. Terminal Services Architecture



**Table 2. Advantages of Terminal Services**

<b>Enhanced Security</b>	<ul style="list-style-type: none"> <li>With the OS, applications, and data locked down in the data center, this model has lower risk of a security breach or data loss via the client than client-based models.</li> </ul>
<b>Simplified Manageability</b>	<ul style="list-style-type: none"> <li>Terminal Services is a mature and well-understood technology.</li> <li>Application and data management are centralized, allowing simpler administration and more reliable backup.</li> <li>Software image management, validation, and support are simplified. Driver and dynamic link library (DLL) conflicts are reduced. Adding, moving, and changing users is simple.</li> </ul>
<b>Lower Cost of Incremental Software Deployment</b>	<ul style="list-style-type: none"> <li>Most enterprises already have some Terminal Services, so adding new applications may not require much in the way of new infrastructure or software.</li> <li>More users can be supported per terminal server than other server-based models, such as VDI.</li> </ul>
<b>Remote Access</b>	<ul style="list-style-type: none"> <li>Centralized computing allows access from any network-connected client. Users do not need to be at "their" workstation.</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>In the event of a data center or worksite disaster, work can shift to other sites relatively easily assuming redundant servers and data storage are in place.</li> </ul>
<b>Reduced Client Power Consumption</b>	<ul style="list-style-type: none"> <li>With thin client terminals, client power consumption is lower than most desktop PCs. However, total power consumption may remain equivalent due to increase in power consumption in the data center.</li> </ul>

**Table 3. Limitations of Terminal Services**

<b>Performance and Responsiveness</b>	<ul style="list-style-type: none"> <li>With even moderately compute-intensive applications, system responsiveness degrades as the number of users increases. Network bandwidth and loading become major factors in client system performance. Networks must have sufficient peak capacity to support the number of terminal users.</li> <li>Remote Display Protocol (RDP) creates a serious graphics bottleneck. Motion graphics such as video or Adobe Flash* do not perform effectively over RDP. A higher performance, third-party display protocol may be required.</li> <li>Users expect rapid response to mouse clicks and keystrokes. Terminal servers must be physically located close enough to their clients to meet responsiveness requirements.</li> </ul>
<b>Software Compatibility</b>	<ul style="list-style-type: none"> <li>Not all software runs under Terminal Services. Also, Voice over Internet Protocol (VoIP), streaming media, and compute- or graphics-intensive applications may be poorly suited to server-based computing.</li> </ul>
<b>Lack of Mobility</b>	<ul style="list-style-type: none"> <li>Application delivery via Terminal Services requires a persistent network connection with adequate bandwidth. Wireless laptops or tablets can be used with Terminal Services, but the session ends if the device disconnects from the network.</li> </ul>
<b>Cost of New User Deployment</b>	<ul style="list-style-type: none"> <li>New deployment to numerous users may involve significant expense, such as new servers, software, network storage, and network infrastructure.</li> <li>Acquisition cost of thin client terminals is similar to many desktop PCs. Initial cost of Terminal Services is about the same as VDI, but higher than streaming or well-managed PCs<sup>1</sup>.</li> </ul>
<b>Single Point of Failure</b>	<ul style="list-style-type: none"> <li>Loss of data center or network function takes all users offline unless backed by a redundant system.</li> </ul>
<b>Lower User Satisfaction</b>	<ul style="list-style-type: none"> <li>Users are conditioned to the PC experience in terms of performance, customization, flexibility, and mobility. Inappropriate or overly strict implementation of Terminal Services can result in high user dissatisfaction and complaints.</li> </ul>

## Virtual Desktop Infrastructure (VDI)

### Summary

Virtual Desktop Infrastructure (VDI) is a newer model that is attracting a great deal of industry buzz, although it's not yet as widely implemented or well understood as Terminal Services.

VDI is designed to offer the responsive and customizable user experience of rich distributed computing along with the management and security advantages of server-based models. It promises centralized management of the entire desktop image.

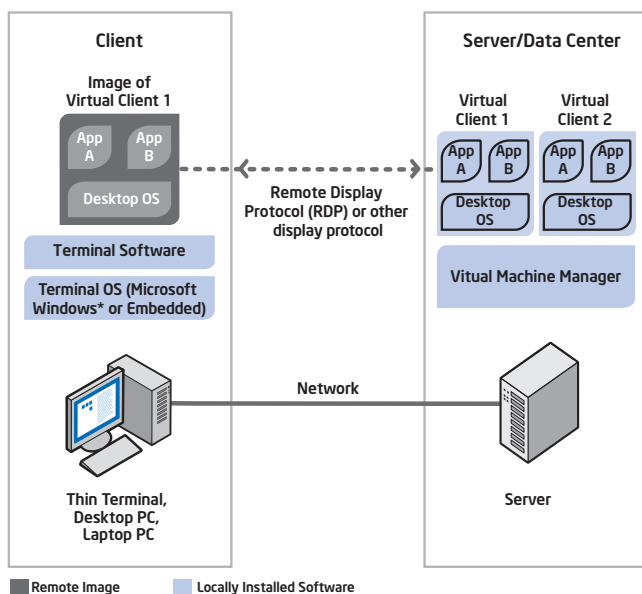
As with other server-based models, performance and responsiveness vary depending on number of users, physical distance, and type of application. Video, Adobe Flash,\* Voice over Internet Protocol (VoIP), and other compute- or graphics-rich applications are not well suited to this model without additional media acceleration. Also, VDI requires a persistent network connection, so it's not appropriate where off-network mobility is required.

New deployments of VDI to large numbers of users may be costly. All client computation, graphics, and memory resources must be built into the data center, and the storage system must accommodate OS, applications, and data for each user. IT departments need to weigh these infrastructure costs against potential TCO savings in manageability.

### How it Works

As with Terminal Services, all computation and storage are centralized, with application images pushed over the network to the client via Remote Display Protocol (RDP) or other display protocols. The major difference is that VDI can offer each user their own complete virtual machine and customized desktop, including the OS, applications, and settings.

Figure 2. VDI Architecture



**Table 4.** Advantages of VDI

Enhanced Security	<ul style="list-style-type: none"> <li>With the OS, applications and data locked down in the data center, this model has lower risk of a security breach or data loss via the client than client-based models.</li> </ul>
Simplified Manageability	<ul style="list-style-type: none"> <li>Application and data management are centralized, allowing simpler administration and more reliable backup.</li> <li>Software image management, validation, and support are simplified. Driver and dynamic link library (DLL) conflicts are reduced.</li> </ul>
User Customization	<ul style="list-style-type: none"> <li>Each user owns a full virtual machine on the server, allowing PC-like personalization of preferences and settings.</li> </ul>
Remote Access	<ul style="list-style-type: none"> <li>Centralized computing allows access from any network-connected client. Users do not need to be at "their" workstation.</li> </ul>
Disaster Recovery and Business Continuity	<ul style="list-style-type: none"> <li>In the event of a data center or worksite disaster, work can shift to other sites relatively easily assuming redundant servers and data storage are in place.</li> </ul>
Reduced Client Power Consumption	<ul style="list-style-type: none"> <li>With thin client terminals, client power consumption is lower than most desktop PCs. However, total power consumption may remain equivalent due to increase in power consumption in the data center.</li> </ul>

**Table 5.** Limitations of VDI

Performance and Responsiveness	<ul style="list-style-type: none"> <li>With even moderately compute-intensive applications, system responsiveness degrades as the number of users increases. Network bandwidth and loading become major factors in client system performance. Networks must have sufficient peak capacity to support the number of VDI users.</li> <li>Remote Display Protocol (RDP) creates a serious graphics bottleneck. Motion graphics such as video or Adobe Flash* do not perform effectively over RDP. A higher performance, third-party display protocol may be required.</li> <li>Users expect rapid response to mouse clicks and keystrokes. Terminal servers must be physically located close enough to their clients to meet responsiveness requirements.</li> </ul>
Manageability	<ul style="list-style-type: none"> <li>Although software images don't reside on the client, IT must still manage, update, and patch all the virtual desktop images now stored in the data center.</li> </ul>
Software and Device Compatibility	<ul style="list-style-type: none"> <li>Applications, devices and their associated drivers are intermediated with the hardware by the VMM. Devices that are highly specialized or software that requires direct hardware interaction may experience compatibility issues.</li> </ul>
Lack of Mobility	<ul style="list-style-type: none"> <li>Desktop delivery via VDI requires a persistent network connection with adequate bandwidth. Wireless laptops or tablets can be used with VDI, but the application session ends if the device disconnects from the network.</li> </ul>
Cost of New User Deployment	<ul style="list-style-type: none"> <li>New deployments of VDI to large numbers of users may involve significant expense. All client computation, graphics, and memory resources must be built into the data center, and the storage system must accommodate OS, application, and data for each user. New hardware, networking, and building space may be required.</li> <li>The number of users supported per VDI server is lower than other server-based models.</li> <li>Acquisition cost of thin client terminals is similar to many desktop PCs.</li> </ul>
Single Point of Failure	<ul style="list-style-type: none"> <li>Loss of data center or network function takes all users offline unless backed by a redundant system.</li> </ul>
Lower User Satisfaction	<ul style="list-style-type: none"> <li>Users are conditioned to the PC experience in terms of performance, flexibility, and mobility. Inappropriate or overly strict implementation of VDI can result in high user dissatisfaction and complaints.</li> </ul>

## Blade PCs

### Summary

Blade PCs bring client computing into a central location, promising higher manageability and security than rich distributed computing through restricted physical access, software imaging policies, and limits on user activities.

If each user is assigned a single PC blade (one-to-one), the model most closely resembles rich distributed computing, except it can only be used in a fixed location and users must be constantly connected to the network. If individual PC blades service multiple users simultaneously (one-to-many), this model more closely resembles VDI.

The proprietary nature of blade computing makes it relatively costly to implement and requires IT organizations to make a long-term commitment to a vendor-specific architecture. Switching vendors or compute models may require a complete “forklift upgrade.” In addition, one-to-many blade PCs may demonstrate some of the same drawbacks as VDI: performance issues and high infrastructure costs.

### How it Works

Blade PCs repartition the PC, leaving basic display, keyboard, and mouse functions on the client, and putting the processor, chipset, and graphics silicon on a small card (blade) mounted in a rack on a central unit. OS, application, and data storage are centralized in a storage array.

Unlike server blades, PC blades are built from standard desktop or mobile processors and chipsets. The central unit, which supports many individual blades, is secured in a data center or other IT-controlled space. In some cases, remote display and I/O is handled by dedicated, proprietary connections rather than using RDP over the data network.

Blade PC vendors initially targeted a user-to-blade ratio of one-to-one, where each user was dynamically assigned a blade and had exclusive use of it. As blade solutions and virtualization software have advanced, most vendors are now offering one-to-many capabilities.

Figure 3. One-to-One Blade Architecture

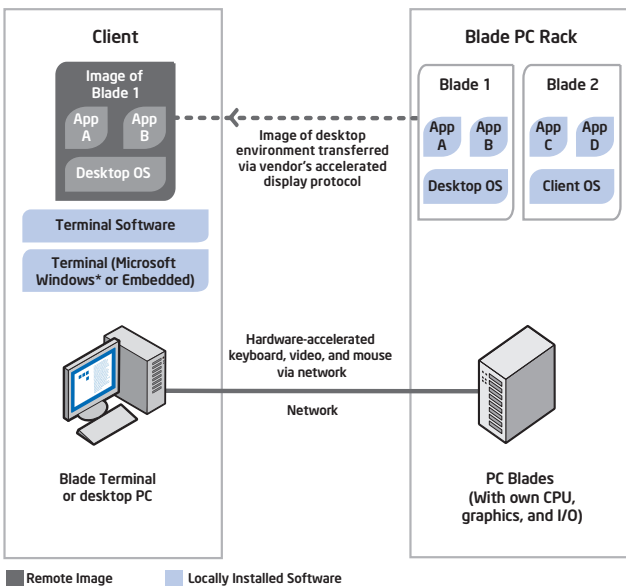
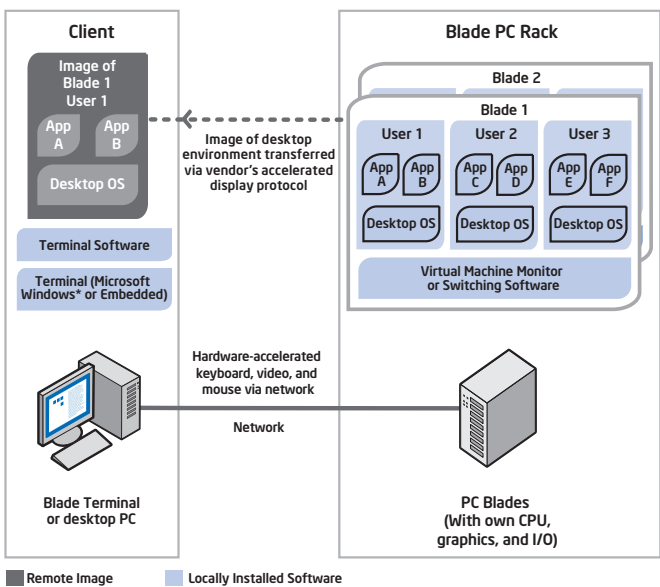


Figure 4. One-to-Many Blade Architecture





**Table 6. Advantages of Blade PCs**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ With the OS, applications and data locked down in the data center, this model has lower risk of a security breach or data loss via the client than client-based models.</li> </ul>
<b>Manageability</b>	<ul style="list-style-type: none"> <li>▪ OS, application, and data management are centralized, allowing simpler administration and more reliable backup.</li> <li>▪ PC blades offer a uniform hardware platform, which simplifies validation, image management, and support.</li> <li>▪ Users can be dynamically assigned any PC blade that is available. Adding, moving, and changing users is relatively simple.</li> </ul>
<b>User Customization</b>	<ul style="list-style-type: none"> <li>▪ Each user can have a unique OS and application image, allowing PC-like personalization of preferences and settings.</li> </ul>
<b>Remote Networked Access</b>	<ul style="list-style-type: none"> <li>▪ Centralized architecture allows users access from any blade-connected client, or in some cases, any Internet-connected client. Users do not need to be at “their” workstation for access.</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>▪ In the event of a blade rack or worksite disaster, work can shift to other sites where redundant infrastructure is available.</li> <li>▪ A user can easily be migrated to another blade if their assigned blade fails.</li> </ul>
<b>Lower Client Power Consumption</b>	<ul style="list-style-type: none"> <li>▪ Power consumption, heat, and fan noise at the client device are lower than most desktop PCs, although total power consumption of clients, blades, and associated storage may be comparable.</li> </ul>

**Table 7. Limitations of Blade PCs**

<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ In one-to-many deployment, application performance may degrade depending on the number of users and their workloads.</li> </ul>
<b>Manageability</b>	<ul style="list-style-type: none"> <li>▪ Although software images don’t reside on the client, IT still must manage, update, and patch all of the centralized desktop images stored in the data center.</li> </ul>
<b>Vendor Lock-in</b>	<ul style="list-style-type: none"> <li>▪ Blade PCs are not standardized, and each vendor has a proprietary implementation. Once IT has selected a vendor’s blade architecture, the costs of switching may be extremely high. Switching vendors or compute models requires a complete hardware upgrade in most cases.</li> <li>▪ Available management tools may be limited, and IT may be dependent on their blade vendor’s tools and development schedule.</li> </ul>
<b>Lack of Mobility</b>	<ul style="list-style-type: none"> <li>▪ No mobile option exists for blade PCs. They are only suitable for users with a persistent network connection.</li> </ul>
<b>Higher Cost Per User</b>	<ul style="list-style-type: none"> <li>▪ Due to their non-standard architecture, blade PC acquisition costs per user are higher than other models.</li> <li>▪ The storage system must accommodate the OS, application, and data for new users. New hardware, networking, and building space may be required.</li> </ul>
<b>Single Point of Failure</b>	<ul style="list-style-type: none"> <li>▪ Loss of a blade server, network access, or data center takes all users offline unless backed by a redundant system.</li> </ul>

## OS Image Streaming or Remote OS Boot

### Summary

Like VDI, OS Image Streaming and Remote OS Boot promise centralized management of the entire desktop image. These models offer users the responsiveness and performance advantages of local execution, while IT gets the manageability and security benefits of centralization without the larger infrastructure build-out often required by VDI.

Currently, no vendor offers OS Image Streaming or Remote OS Boot product with local caching, so these models cannot be used where mobility is required. Because data resides on the client during execution, it may not be suitable for applications that require the highest data security.

### How it Works

At startup, the client is essentially “bare metal,” with no OS Image installed locally. The OS Image is streamed to the client over the network, where it executes locally using the client’s own CPU and graphics. Application data is stored in a data center. The client is usually a PC with no hard drive, which uses RAM exclusively.

With streaming technology, the OS Image software does not stream to the clients in the same form as it comes from the software vendor. The software first goes through a preparation process, where it is divided up into prioritized blocks and placed in a specific order for streaming to the client. The basic launch and initiation software go first, followed by high-demand services and capabilities. These optimizations allow the OS to launch and begin operations, even before all the code is streamed to the client. In order to reduce network traffic, some less frequently used capabilities may remain in the data center until requested.

Remote OS Boot is similar to OS Image Streaming in that it delivers a complete OS Image to a “stateless” PC whose hard drive has been deactivated or removed. Unlike OS Image Streaming, Remote OS Boot clients boot directly from the storage area network (SAN) device; The PC treats the SAN just like its local hard drive. The OS image is unmodified from the “gold” image that would be used on a local disk.

Figure 5. OS Image Streaming Architecture

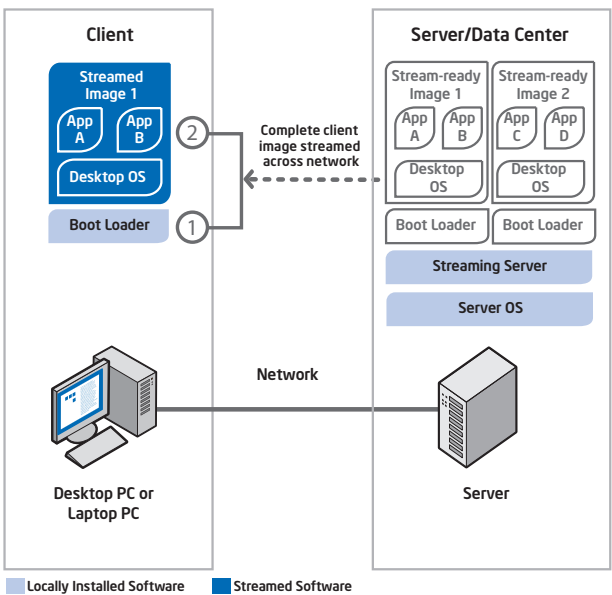
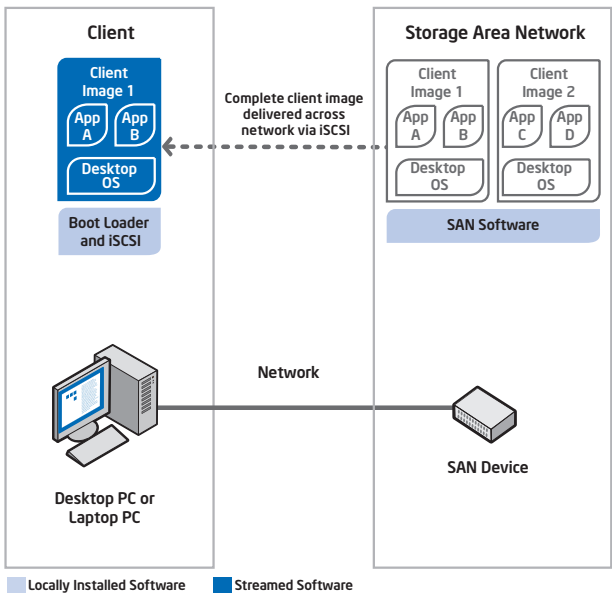


Figure 6. Remote OS Boot Architecture



**Table 8. Advantages of OS Image Streaming and Remote OS Boot**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ Critical application data is stored in the data center.</li> <li>▪ Any application or OS corruption is eliminated, and patches automatically applied, when the system is reimaged at startup.</li> </ul>
<b>Manageability</b>	<ul style="list-style-type: none"> <li>▪ OS, application, and data management are centralized, allowing simpler administration, easy software migration, and more reliable backup.</li> <li>▪ OS Image Streaming and Remote OS Boot significantly reduce the image management and the desktop software support issues found with locally-installed software.</li> <li>▪ Software licensing can be centrally managed. Streaming provides better insight into actual application usage, enabling greater licensing optimization.</li> <li>▪ Stateless clients make adding, moving, and changing users very simple.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Performance is virtually identical to traditional locally installed applications. Compute- and graphics-intensive applications, as well as video, Adobe Flash,* and streaming media, all perform well.</li> <li>▪ After initial boot-up, network load is lower than server-based solutions, because Remote Display Protocol (RDP) is not used to push application screen images over the network.</li> </ul>
<b>Infrastructure Cost Savings</b>	<ul style="list-style-type: none"> <li>▪ Less server and network infrastructure is required compared to server-based models.</li> <li>▪ OS Image Streaming technology has the lowest TCO of all centralized compute models!</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>▪ In the event of a data center or worksite disaster, work can shift to other sites relatively easily assuming redundant servers or a data center is in place.</li> </ul>

**Table 9. Limitations of OS Image Streaming and Remote OS Boot**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ Data and applications reside on the client at runtime and are therefore somewhat more susceptible to client-side attacks than server-side models.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Booting a streamed OS Image over the network can be somewhat slower than booting locally.</li> <li>▪ Streaming download speeds are affected by physical distance from the server, network load, and number of users.</li> </ul>
<b>Software Compatibility and Implementation Issues</b>	<ul style="list-style-type: none"> <li>▪ Due to its internal architecture, some legacy or custom application software cannot successfully complete the preparation process and thus cannot be delivered via OS Image Streaming or Remote OS Boot.</li> <li>▪ With OS Image Streaming, the initial setup and debug of the software preparation process can be time- and labor-intensive.</li> </ul>
<b>Lack of Mobility</b>	<ul style="list-style-type: none"> <li>▪ As of this writing, no commercially available product allows off-network or mobile use of OS Image Streaming or Remote OS Boot models.</li> </ul>

## Application Streaming or Application Virtualization

### Summary

Like Terminal Services, Application Streaming enables centralized application management, but does so without sacrificing responsiveness or performance. Data can be stored locally or centrally, depending on enterprise policy. Plus, streamed applications can be cached for off-network use, making it a centrally-managed model that supports mobile computing options.

Because data resides on the client during execution, this model may not be suitable for applications that require the highest levels of security. Application Virtualization can also introduce limitations in how applications interact.

### How it Works

The client OS is locally installed, but applications are streamed on demand from the server to the client, where they are executed locally.

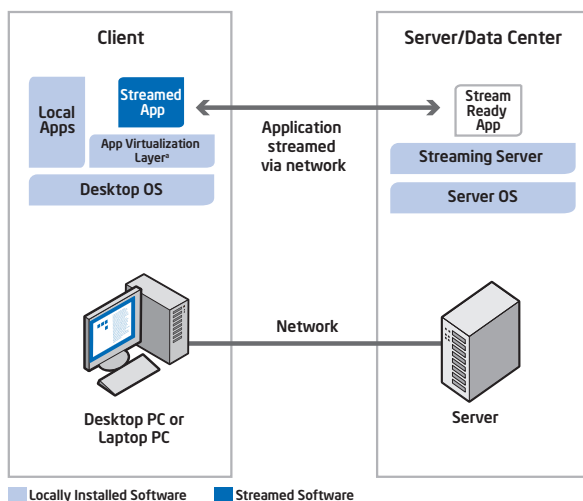
Although the terms “streaming” and “application virtualization” are often used interchangeably, they are not the same thing. Streaming refers to the delivery model of sending the software over the network for execution on the client. Streamed software can be installed in the client OS locally or, in most cases, it is virtualized.

With Application Virtualization, streamed software runs on an abstraction layer and does not install in the OS registry or system files. This simplifies the interactions between the streamed application and the OS, significantly reducing software conflicts and image management problems.

However, some virtualized applications may not interact with other applications as they do when they are both locally installed – for example, cut and paste may not work. This issue can be mitigated by streaming and virtualizing related applications in bundles, or by using the most recent releases of virtualization software, which have addressed most of these issues.

Unlike other centralized computing models, streamed applications can be cached on a laptop and taken off the network. When the laptop reconnects with the network, the application can resynchronize with the server, check licensing and patch information, and download application data to the data center.

**Figure 7.** Application Streaming or Application Virtualization Architecture



\*Streaming ISV's agent. Degree of virtualization and isolation varies by vendor and policy.

**Table 10. Advantages of Application Streaming or Application Virtualization**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ Critical application data can be stored in the protected data center based on policy. Local Storage can be disallowed.</li> <li>▪ Application corruption is eliminated, and patches automatically applied, when the application is reloaded at startup.</li> <li>▪ Virtualized applications can be isolated from each other, limiting data exposure to other applications or the OS.</li> </ul>
<b>Manageability</b>	<ul style="list-style-type: none"> <li>▪ Applications, and in some cases, data are centralized, allowing simpler administration, easier software migration, and more reliable backup.</li> <li>▪ Application virtualization significantly reduces the image management as well as the desktop software support issues found with locally-installed software.</li> <li>▪ Application virtualization may enable legacy applications to run on a newer OS, even if the application has compatibility problems when locally installed.</li> <li>▪ Software licensing can be centrally managed. Streaming provides better insight into actual application usage, enabling greater licensing optimization.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Performance is virtually identical to traditional locally installed applications. Compute- and graphics-intensive applications, as well as video, Adobe Flash,* and streaming media, all perform well.</li> <li>▪ Caching options can be set to accelerate initial startup and application launch without storing application data locally.</li> <li>▪ Network load is lower than server-based solutions, because application screen images do not need to be pushed over the network using Remote Display Protocol (RDP) or other protocol.</li> <li>▪ Streaming only the applications reduces the network load compared to streaming the entire OS Image.</li> <li>▪ User experience is the same as local rich client.</li> </ul>
<b>Infrastructure Cost Savings</b>	<ul style="list-style-type: none"> <li>▪ Less server and network infrastructure is required compared to server-based models.</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>▪ Users can continue to work on local clients with cached applications even if the network or data center is offline.</li> </ul>
<b>Mobility</b>	<ul style="list-style-type: none"> <li>▪ Unlike other server-based computing models, streamed applications can be cached for off-network use on mobile clients.</li> </ul>

**Table 11. Limitations of Application Streaming or Application Virtualization**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ Data and applications reside on the client at runtime and are therefore somewhat more susceptible to client-side attacks or theft than server-based models.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Streaming download speeds are affected by physical distance from the server, network load, and number of users.</li> <li>▪ Virtualization may limit interactions between applications (for example, no cut and paste between applications).</li> </ul>
<b>Software Compatibility and Implementation Issues</b>	<ul style="list-style-type: none"> <li>▪ Due to its internal architecture, some legacy or custom application software cannot successfully complete the preparation process and thus cannot be delivered via Application Streaming.</li> <li>▪ With application virtualization, the initial setup and debug of the software preparation process can be time- and labor-intensive.</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>▪ Although application data is stored centrally, getting users working again in a new site is relatively more complex than other models due to locally installed OS and/or local data.</li> </ul>

## Virtual Containers

### Summary

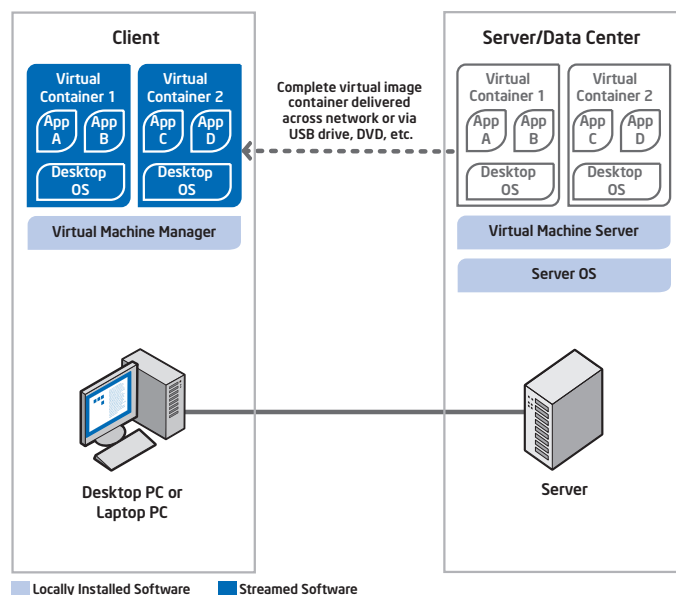
“Virtual Containers” is a new, but rapidly evolving, compute model. It provides centralized management of the complete OS and application image. Unlike locally installed or steamed OSs, the Virtual Container is abstracted away from the client hardware by a virtual machine manager (VMM). The presence of the VMM eases the validation burden for IT, since various platform hardware differences are hidden from the OS in the container. It also enables new usage models, such as managed contactor images and “bring-your-own-PC” programs. Since it uses client-side execution, a large data center build-out is not required, and where server-based execution makes sense, the same Virtual Containers could be run on a VDI server. For users, Virtual Containers offer responsive local execution and off-network mobility.

Virtual Containers can be general-purpose user environments, such as a standard Windows OS and productivity applications, or they can be purpose-built, single function “virtual appliances” that provide services like compliance monitoring or highly secure applications.

### How it Works

In this model, virtual machine images, including the OS and applications, are created and managed centrally by IT. But instead of running the virtual machine on the server (the VDI model), the virtual machine is streamed to the client for local execution on a client-based VMM. Since execution is on the client, even compute- or graphics-intensive applications are responsive, and users can enjoy off-network mobility. There are several technical issues that the industry is striving to solve, such as graphics virtualization, wireless and power management.

Figure 8. Virtual Containers Architecture



**Table 11. Advantages of Virtual Containers**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ With Virtual Containers, virus and intrusion threats are contained and applications requiring high security can be easily isolated in separate containers.</li> <li>▪ Security policies can be set for the needs of the Virtual Container, allowing IT more control over lock-down, access policy, data storage and authorization/revocation rights.</li> <li>▪ Purpose-built virtual appliances can provide valuable security services from outside the user's environment.</li> </ul>
<b>Manageability</b>	<ul style="list-style-type: none"> <li>▪ OS images, application and data management can be centralized, allowing simpler administration, easy software migrations and more reliable back-up.</li> <li>▪ Since Virtual Containers run on a VMM, not the physical hardware, image validation only needs to be done against the VMM, not on many unique hardware configurations.</li> <li>▪ Virtual Containers are highly portable. Installation may be as simple as streaming a file to the client or inserting a USB flash drive.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Performance is virtually identical to traditional locally installed applications, compute- and graphics-intensive applications, as well as video, Adobe Flash,* and streamed media, all perform well.</li> <li>▪ VMM efficiency and performance will be a factor.</li> </ul>
<b>Mobility</b>	<ul style="list-style-type: none"> <li>▪ Unlike other server-based computing models, Virtual Containers can be cached for off-network use on mobile clients.</li> <li>▪ Users could carry their Virtual Container image on a USB drive and run it on any PC (home, office, partner site, etc.).</li> </ul>
<b>Infrastructure Cost Savings</b>	<ul style="list-style-type: none"> <li>▪ Less server and network infrastructure is required compared to server-based models.</li> </ul>
<b>Disaster Recovery and Business Continuity</b>	<ul style="list-style-type: none"> <li>▪ Users can continue to work on local clients using cached Virtual Containers even if the network or data center is offline.</li> </ul>

**Table 12. Limitations of Virtual Containers**

<b>Security</b>	<ul style="list-style-type: none"> <li>▪ At runtime, data and applications are resident on the client, and therefore more susceptible to client-side attacks or theft than server-side models.</li> <li>▪ VMM is a new layer in the software stack that must also be protected.</li> </ul>
<b>Performance</b>	<ul style="list-style-type: none"> <li>▪ Running multiple virtual machines on a VMM may cause performance degradation. VMM efficiency and hardware-assistance will be key to achieving "near native" application performance.</li> <li>▪ Virtualization may limit interactions between applications. This may be by design in some cases, but may be an unintended consequence in others.</li> </ul>
<b>Maturity</b>	<ul style="list-style-type: none"> <li>▪ The Virtual Containers model is relatively new. The technology, management tools, and IT processes are not fully mature.</li> </ul>
<b>Industry-Wide Technical Challenges</b>	<ul style="list-style-type: none"> <li>▪ The industry is striving to solve several technical challenges before the full potential of Virtual Containers can be realized. This includes virtualization of graphics, wireless, power management, docking stations and peripherals.</li> </ul>

## Factors to Consider

Choosing the appropriate compute model requires balancing several interrelated factors. Excessive focus on only one of these factors is likely to result in a less than optimal solution. For example, focusing exclusively on IT requirements may result in an easy-to-manage system that users can't stand.

<b>IT Requirements</b> Standard IT concerns, such as security, manageability, business continuity, etc	<ul style="list-style-type: none"><li>▪ Security</li><li>▪ Image management</li><li>▪ License management</li><li>▪ Support structure</li><li>▪ Disaster recovery</li></ul>
<b>Infrastructure</b> The hardware, connectivity, and bandwidth available to deliver the compute model. If the infrastructure is not in place to support a particular compute model, either it needs to be purchased and installed, or another model must be chosen.	<ul style="list-style-type: none"><li>▪ Servers</li><li>▪ Storage</li><li>▪ Data center space, power, and cooling</li><li>▪ Network bandwidth</li><li>▪ Budget priorities</li></ul>
<b>User Experience</b> The workflow needs of the system's users, including the need for mobility and performance. In many cases, either by custom or policy, users expect that the device may be used for some personal tasks, which may result in personally identifiable information being on the system. These issues of "ownership" and privacy should be considered.	<ul style="list-style-type: none"><li>▪ Mobility</li><li>▪ Responsiveness</li><li>▪ Customization</li><li>▪ Connectivity</li><li>▪ "Ownership" and privacy</li></ul>
<b>Application Workload</b> The compute and graphics demands of the applications the user will be running. Some applications may be intolerant of network delays, such as VoIP or streamed video. Headroom for future application growth should also be considered.	<ul style="list-style-type: none"><li>▪ Compute load</li><li>▪ Graphics load</li><li>▪ Delay sensitivity (e.g., video, motion graphics, VoIP)</li><li>▪ Headroom</li></ul>

## Conclusion

As software delivery models have evolved, IT departments have more choices than ever. Choosing the appropriate compute model requires a balance of many factors. Finding a solution that meets all relevant user and IT needs is likely to involve a mix of delivery models, even within a single user segment.

Desktop PCs with the Intel Core 2 processor with vPro technology and Intel Centrino Pro with vPro technology for notebooks offer highly manageable, secure platforms on which to deliver all compute models.

## Additional Resources

Principled Technologies offers free white papers, which are a great source of third-party quantified data to support the points in this document. They are available on Principled Technologies' web site [www.principledtechnologies.com](http://www.principledtechnologies.com).

- White Paper: Total Cost of Ownership for Various Computing Models [www.principledtechnologies.com/Clients/Reports/Intel/CompModelsTCO1107.pdf](http://www.principledtechnologies.com/Clients/Reports/Intel/CompModelsTCO1107.pdf)
- Spreadsheet: Total Cost of Ownership Calculator [www.principledtechnologies.com/Clients/Reports/Intel/ComputeModelTCOCalc1107.xls](http://www.principledtechnologies.com/Clients/Reports/Intel/ComputeModelTCOCalc1107.xls)

<sup>1</sup> Principled Technologies, "Total Cost of Ownership for Various Computing Models." See Additional Resources for reference.

\*Other names and brands may be claimed as the property of others.

Copyright © 2009 Intel Corporation. All rights reserved. Intel, the Intel logo, Centrino, Intel Core, and Intel vPro are trademarks of Intel Corporation in the U.S. and other countries.

Printed in USA

0109/WW/OCG/XX/PDF

 Please Recycle

319175-002US

